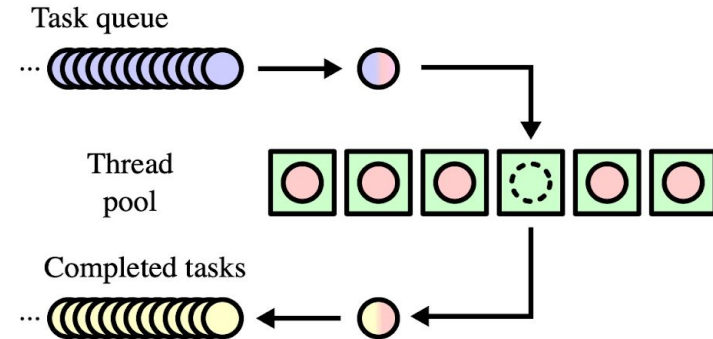



# ThreadPool

Harshit Roy  
UT Austin



# BOOST ASIO

- **Threads** fetch tasks from a **central queue**
- **Scheduling policies:**
  - **FIFO** (First-In, First-Out)
  - **Round-robin** (load-balanced)
- Uses **epoll** to monitor I/O readiness without blocking threads
- One thread can wait on Linux **epoll** fd, thus limiting scalability
- **No dynamic decisions** - least overhead




# Intel TBB

**No central task queue** — improves scalability & reduces contention

## Cilk-style Work Stealing

- Work: LIFO (Last-In, First-Out)
  - Great for locality and cache reuse
- Steal: FIFO (First-In, First-Out)
  - Load Balancing, oldest tasks are most independent
- LIFO + FIFO
  - A sweet spot

**Dynamic Decision** - results in overhead




# BS ThreadPool

**No central task queue** — improves scalability & reduces contention

**Dynamic Scheduling** - high overhead

Similar to Intel TBB for Compute and IO tasks

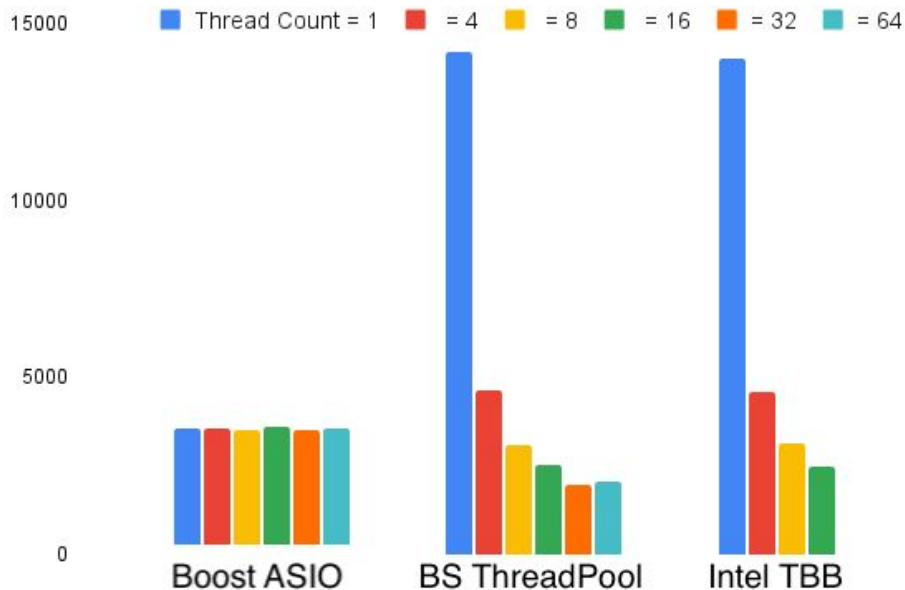
Inefficient in handling mixed workloads (compute + I/O)



# I/O-Intensive Workload: File Read/Write Throughput


Creates 50 tasks (NUM\_TASKS = 50)

- Each task:
  - Writes a 100 MB file to disk filled with 'A'
  - Then reads the same file back into memory



## Boost ASIO


- Designed for async IO tasks
- Scaling is limited



# Compute-Intensive Workload


Creates three square matrices:

- A and B are filled with random integers from 1 to 10
- C is initialized to all 0s, to store the result of  $A \times B$



## Boost ASIO

- Not designed for CPU-intensive tasks
- Scaling is limited




# Mixed Workload ( Compute + IO Workload )

Splits data into chunks of 1 million integers.


Each chunk is independently: extracted, sorted, saved to a separate chunk file

These can later be **merged** later



## BS ThreadPool


- Not designed for mixed workload



# Thread Pool Overhead

Task scheduling overhead

Time taken to assign no-op tasks to threads.




BS ThreadPool

- High overhead

Boost ASIO

- Least overhead






# Thread Scheduling

CPU-intensive and I/O Intensive tasks are submitted concurrently.

None of the thread pools support non-blocking I/O

E.g. No of threads = 4, 5 I/O and CPU Task submitted concurrently

I/O Task 0	started at 399985 ms,	ended at 411874 ms,	duration: 11.88 sec
CPU Task 1	started at 399985 ms,	ended at 418241 ms,	duration: 18.25 sec
CPU Task 2	started at 399985 ms,	ended at 418291 ms,	duration: 18.30 sec
CPU Task 0	started at 399985 ms,	ended at 418368 ms,	duration: 18.38 sec
I/O Task 1	started at 411874 ms,	ended at 424708 ms,	duration: 12.83 sec
I/O Task 3	started at 418370 ms,	ended at 430772 ms,	duration: 12.40 sec
I/O Task 2	started at 418244 ms,	ended at 430896 ms,	duration: 12.65 sec
CPU Task 3	started at 418293 ms,	ended at 436642 ms,	duration: 18.34 sec
I/O Task 4	started at 430772 ms,	ended at 442698 ms,	duration: 11.92 sec
CPU Task 4	started at 424708 ms,	ended at 442922 ms,	duration: 18.21 sec



# Questions?

[https://github.com/royharshit/threadpool\\_benchmark](https://github.com/royharshit/threadpool_benchmark)

